

The Longest-Answer Heuristic: Characterizing and Mitigating Answer-Length Bias in LLM-Generated Multiple-Choice Questions as a Threat to Assessment Integrity

Steven*

Department of Artificial Intelligence
Kyungdong University Global Campus

Abstract

Large language models (LLMs) are increasingly used by educators and students to author multiple-choice questions (MCQs) at scale. We investigate a specific and underexamined failure mode: a systematic correlation between the *length* of an answer option and its probability of being keyed correct. When present, this correlation reintroduces a classical item-writing flaw—“the longest option is correct”—at machine scale, furnishing test-takers with a construct-irrelevant “pick-the-longest” heuristic that inflates scores without measuring knowledge. We make four contributions. First, we formalize *answer-length leakage* and define a reusable, model-agnostic measurement suite: Longest-Answer Accuracy (LAA), an *Exploitability* index E that normalizes a guessing strategy’s advantage over chance, a verbosity differential δ , and a length-only leakage AUC. Second, we release an open, provider-agnostic data-collection harness that elicits MCQs from six widely used assistants under an identical, length-neutral prompt, so that the bias can be *measured* rather than assumed. Third, we contribute a generation-time mitigation toolkit—an auditing gate, a length-balanced rewrite planner, length-matched adversarial distractors, and form-level key-length balancing—with a measured before/after protocol. Fourth, we provide a fully reproducible analysis pipeline. Because we did not have sanctioned API access to all six commercial systems at submission time, the quantitative results reported here are generated from a clearly labeled *synthetic* corpus that exercises the entire pipeline end to end; on that corpus the toolkit reduces exploitability from $E=+0.36$ to near parity ($E=-0.06$) and cuts leaky items from 71% to 3%. We position the measurement suite and mitigations as deployable safeguards and provide a preregistered protocol for the real-model and human-subject studies needed to confirm the effect in the wild.

Keywords: large language models; multiple-choice ques-

tions; automatic question generation; assessment validity; item-writing flaws; construct-irrelevant variance; test-wiseness; verbosity bias; educational integrity; mitigation.

1 Introduction

Automatic question generation (AQG) with large language models (LLMs) has moved from research prototype to everyday classroom tool. Instructors use chat assistants to draft quizzes in seconds; students use the same assistants to generate practice items; publishers and learning platforms embed them in content pipelines. The appeal is obvious—authoring high-quality multiple-choice questions (MCQs) by hand is slow and expensive [1]. But a question that is cheap to produce is not necessarily *valid*. An MCQ is valid only insofar as a correct response reflects the intended knowledge or skill and nothing else. When a surface feature of the item—unrelated to its content—predicts the keyed answer, the item leaks. Test-takers who learn to read that feature can score above chance without knowing the material, and the score loses its meaning [2].

This paper concerns one such surface feature: **the length of the answer options**. A long-standing rule of professional item writing is that the correct option should not be conspicuously longer than the distractors, precisely because writers tend—often unconsciously—to over-qualify the key with hedges, conditions, and exceptions, making it both longer and more “textbook-correct” in appearance [3, 4]. Skilled test-takers exploit this: choosing the longest option is a documented component of *test-wiseness*, the ability to use cues in an item’s form rather than its content [5]. The concern we raise is that LLMs, trained on enormous corpora of human-written text and optimized with human feedback that can reward thoroughness and verbosity [6, 7], may reproduce or even amplify this tendency—generating keys that are systematically longer than the distractors they sit beside. If so, every quiz produced this way ships with a built-in cheat code, and the more an institution leans on LLM-authored assessment, the

*Department of Artificial Intelligence, Kyungdong University Global Campus, Chuncheon, Republic of Korea. Correspondence: author@kduglobal.ac.kr.

more pervasive the leak becomes.

Why this matters now. Three trends compound the risk. (i) *Scale*: a single prompt can yield hundreds of items, so a small per-item bias becomes a large aggregate exposure. (ii) *Automation bias*: users tend to trust machine-generated artifacts and review them less critically, so flawed items are less likely to be caught in editing [8]. (iii) *Homogeneity*: when many instructors use the same few models, the same exploitable pattern recurs across courses and institutions, making the heuristic easy for students to learn and transfer. A bias that an individual human writer might exhibit idiosyncratically becomes, under LLMs, a systematic and *predictable* property of the tool.

Objectives. We set out to (1) make the phenomenon *measurable* with precise, reusable metrics; (2) make it *collectable* with an auditable, provider-agnostic harness that elicits items under a length-neutral prompt; (3) make it *fixable* at generation time with lightweight, deployable mitigations; and (4) make the whole pipeline *reproducible*. We deliberately treat the prevalence of the bias as an empirical question to be answered by measurement, not as a settled fact.

Research questions.

RQ1. Can answer-length leakage in LLM-generated MCQs be defined and quantified with a model-agnostic, exploitability-oriented metric suite?

RQ2. How large is the leakage, and does it vary by model, subject, and difficulty?

RQ3. Does a “pick-the-longest” strategy yield an above-chance score, and how much information does length alone carry about the key?

RQ4. Can generation-time mitigations remove the leak without degrading item quality, and by how much (measured before/after)?

Contributions and scope. Our contributions are the metric suite (§3), the collection harness (§4), the mitigation toolkit (§5), and the reproducible analysis (§7). **An explicit caveat governs the whole paper:** the numbers reported in §7 are produced from a *synthetic* corpus with injected, known bias levels. They exist to validate that the metrics behave correctly and that the mitigations work; they are *not* measurements of any real commercial model and must not be read as such. §10 states this plainly and §11 gives the preregistered protocol for the real-model and human-subject studies required to estimate the effect in deployment.

2 Related Work

2.1 Item-writing flaws and assessment validity

The psychometric literature has catalogued item-writing flaws for decades. Haladyna and Downing [9] introduced a taxonomy of item-writing rules, refined into an evidence-based guideline set by Haladyna et al. [3]; among the rules is the explicit injunction that the correct option should not be made longer than the distractors. Tarrant et al. [4] found such flaws to be common in real examinations and showed they disproportionately affect high-achieving and borderline candidates. Downing [10] demonstrated that flawed items change pass/fail outcomes for a non-trivial fraction of examinees, establishing that these are consequential, not cosmetic, defects. Rodriguez [11] provides meta-analytic evidence on option counts. Underlying all of this is Messick’s framework of *construct-irrelevant variance*: score variance attributable to features other than the target construct [2]. A length cue is a textbook source of such variance, and the ability to use it is the classical notion of *test-wiseness* [5]. Our work imports these constructs into the LLM setting and operationalizes “the longest option is correct” as a measurable, exploitable quantity.

2.2 Automatic question generation

Surveys of AQG document rapid progress and persistent quality-control gaps [1], with distractor generation a particular challenge [12, 13]. The dominant evaluation lens is fluency and answerability; the *statistical validity* of the generated key relative to its distractors—whether the key is distinguishable by a surface feature—has received far less attention. Recent work has begun to flag that AI-assisted authoring can *increase* the rate of item-writing flaws through automation bias [8, 14], and comparative studies of LLM MCQ generators are emerging [15, 16]. We contribute the missing measurement: a direct, cross-model quantification of answer-length leakage as an exploitable cue.

2.3 Biases in LLMs as judges and as answerers

Two adjacent literatures motivate the hypothesis that LLMs would exhibit a length asymmetry. First, when LLMs act as *judges*, they display a *verbosity bias*, preferring longer responses independent of quality [6, 7]. Second, when LLMs act as MCQ *answerers*, they exhibit *selection* and *token bias*—sensitivity to option ordering and to the surface form of option identifiers—making them “not robust” selectors [17, 18]. Both lines concern how models *consume* options. The complementary and largely open question, which we take up, is how models *produce* options: when an LLM writes the

key and its distractors, does it imprint a length signature on the key? If verbosity is rewarded in generation, the key—the option the model is most “committed” to justifying—is exactly where extra length would accrue.

2.4 The gap

Prior work establishes (a) that “longest-correct” is a real, consequential item-writing flaw in human-authored tests; (b) that LLMs are verbose and length-sensitive as judges and answerers; and (c) that AI authoring can raise flaw rates. No prior work, to our knowledge, (i) defines a reusable *exploitability* metric for answer-length leakage, (ii) measures it *across* the major commercial assistants under a common protocol, and (iii) pairs the measurement with a *generation-time* mitigation toolkit evaluated before/after. This paper targets that gap, and is explicit about which parts are delivered as validated tooling versus which require sanctioned data collection to estimate real-world magnitude.

3 Formalizing Answer-Length Leakage

Item model. An MCQ item is a tuple $x = (s, \mathcal{O}, k)$ where s is the stem, $\mathcal{O} = (o_1, \dots, o_m)$ is an ordered list of m options, and $k \in \{1, \dots, m\}$ is the index of the keyed-correct option. Let $\ell(o)$ denote a length measure of option o ; we use character count by default and report word count as a robustness check. Let the random “chance” baseline for a uniform guess be $1/m$.

Longest-Answer Accuracy (LAA). For a corpus $\mathcal{D} = \{x_i\}_{i=1}^N$, define the score a student would obtain by always selecting the longest option (ties broken uniformly at random):

$$\text{LAA}(\mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \frac{\mathbb{I}[k_i \in \arg \max_j \ell(o_{ij})]}{|\arg \max_j \ell(o_{ij})|}. \quad (1)$$

$\text{LAA} = 1/m$ indicates no length signal; $\text{LAA} > 1/m$ indicates the key tends to be the longest option.

Exploitability index E . LAA is not comparable across items with different option counts because the chance baseline $1/m$ differs. We normalize the above-chance advantage to its maximum possible value:

$$E(\mathcal{D}) = \frac{\text{LAA}(\mathcal{D}) - \bar{c}}{1 - \bar{c}}, \quad \bar{c} = \frac{1}{N} \sum_{i=1}^N \frac{1}{m_i}, \quad (2)$$

so $E = 0$ when the strategy does no better than guessing and $E = 1$ when the longest option is *always* the key.

E can be mildly negative if the key is shorter than distractors on average. We adopt E as our primary cross-model, cross-subject comparison statistic because it answers the practically meaningful question: *how much of the way from chance to a perfect cheat does length alone take you?*

Verbosity differential $\bar{\delta}$. To capture magnitude in raw units we define, per item, the gap between the key’s length and the mean distractor length, and average it:

$$\bar{\delta}(\mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \left(\ell(o_{ik_i}) - \frac{1}{m_i - 1} \sum_{j \neq k_i} \ell(o_{ij}) \right). \quad (3)$$

We accompany $\bar{\delta}$ with a standardized effect size (the mean of the per-item differential divided by its standard deviation) and a Wilcoxon signed-rank test of $H_0 : \bar{\delta} = 0$.

Length-only leakage AUC. Finally we ask how much a classifier that sees *only* option lengths can recover the key. For each item we form $(m_i - 1)$ key-vs-distractor pairs and score each option by ℓ ; the AUC is the probability that a randomly chosen key outranks a randomly chosen distractor by length:

$$\text{AUC} = \Pr[\ell(o_{ik_i}) > \ell(o_{ij})], \quad j \neq k_i. \quad (4)$$

$\text{AUC} = 0.5$ is no leakage; $\text{AUC} = 1$ means length perfectly ranks the key first. We complement the AUC with a logistic regression of the binary “is-key” label on standardized option length, reporting the odds ratio per one standard deviation of length, which expresses leakage as a single interpretable multiplier.

All five quantities are implemented in `src/metrics.py` with bootstrap confidence intervals and per-group (model, subject, difficulty) reporting; §7 reports them on the synthetic corpus.

4 Data-Collection Harness

To measure—rather than assume—the bias, we provide `src/collect.py`, a provider-agnostic harness that elicits MCQs from six widely used assistants (Anthropic Claude, OpenAI ChatGPT, Google Gemini, xAI Grok, DeepSeek, and Perplexity) through their respective APIs. The harness embodies three design principles.

Length-neutral elicitation. The generation prompt requests a stem, a fixed number of options, and the index of the correct answer in strict JSON, and *never mentions length, brevity, or option balance*. This is essential: any instruction about length would contaminate the very quantity we measure. The model is free to write options however it “wants,” and the length signature (if any) is therefore an emergent property of the model, not of the prompt.

Balanced, preregisterable design. Requests are issued on a grid of subjects $\times$ difficulty levels with a fixed number of items per cell, so that per-model estimates are not confounded by topic mix. The default grid is five subjects (biology, history, computer science, economics, geography) and three difficulty levels (easy, medium, hard); the cell size is a parameter. The same grid is used for every model, enabling like-for-like comparison and a clean factorial analysis.

Robust parsing and provenance. The harness parses model output defensively (tolerating code fences and minor format drift), validates each item against the schema in §3, records the originating model and request parameters with every item, and writes a JSON-Lines corpus. A `--dry-run` mode produces schema-valid mock items with no network or keys, so the pipeline can be exercised and unit-tested offline; we used this mode to verify end-to-end correctness. The adapter layer is OpenAI-compatible, so additional providers can be added without touching the analysis code.

We stress that running the harness against the live commercial endpoints requires the user’s own API credentials and adherence to each provider’s terms of service; it is the one step we could not execute ourselves (§10).

5 Mitigation Toolkit

We contribute four composable mitigations (`src/mitigate.py`). They are deliberately lightweight so they can run inside an authoring pipeline.

(A) Length auditor (detection gate). The auditor flags any item whose key exceeds the longest distractor by more than a ratio threshold, or whose verbosity differential δ exceeds an absolute threshold. Flagged items are routed to rewriting rather than shipped. This is the minimal, zero-rewrite safeguard and doubles as a continuous monitor.

(B) Length-balanced rewrite planner. For a flagged item, the planner computes a target length band from the distractors and emits a constrained rewrite specification that forbids the key from being the longest option, together with a ready-to-send instruction for an LLM rewriter. A deterministic planner is provided so the step is testable without a model in the loop.

(C) Length-matched adversarial distractors. Rather than shortening the key, this mitigation *lengthens distractors* to match, producing plausible, content-grounded distractors in the key’s length band so that length carries no signal. The toolkit supplies both the prompt builder and a deterministic length-matcher for offline testing.

Algorithm 1: Length Auditor (detection gate)

Input: item $x = (s, \mathcal{O}, k)$; ratio threshold τ_r ; delta threshold τ_δ
Output: FLAG $\in \{\text{LEAKY}, \text{OK}\}$

```

1  $\ell_k \leftarrow \ell(o_k)$ 
2  $\ell_{\max}^- \leftarrow \max_{j \neq k} \ell(o_j)$ 
3  $\delta \leftarrow \ell_k - \frac{1}{m-1} \sum_{j \neq k} \ell(o_j)$ 
4 if  $\ell_k > \tau_r \cdot \ell_{\max}^-$  or  $\delta > \tau_\delta$  then
5   return LEAKY
6 return OK
```

Algorithm 2: Harden Item (generation-time)

Input: item x ; rewriter $\mathcal{R}$ (LLM or deterministic); thresholds τ_r, τ_δ
Output: hardened item x'

```

1 if  $\text{AUDIT}(x, \tau_r, \tau_\delta) = \text{OK}$  then
2   return  $x$ 
3  $\mathcal{B} \leftarrow$  target length band from  $\{\ell(o_j)\}_{j \neq k}$ 
   // Option 1: balance by rewriting the
   // key into band  $\mathcal{B}$ 
4  $o'_k \leftarrow \mathcal{R}_{\text{key}}(o_k, \mathcal{B})$ 
   // Option 2: lengthen distractors into
   // the key’s band
5 for  $j \neq k$  do
6    $o'_j \leftarrow \mathcal{R}_{\text{distractor}}(o_j, \ell(o_k))$ 
7  $x' \leftarrow \text{assemble}(s, \mathcal{O}', k)$ 
8  $x' \leftarrow \text{SHUFFLEOPTIONS}(x')$ 
9 return  $x'$ 
```

(D) Form-level key-length balancing. Across a whole quiz form, the toolkit shuffles option order (removing position cues) and balances the *rank* of the key’s length so that the key is not disproportionately the longest across items. This addresses leakage that is invisible item-by-item but exploitable at the form level.

Algorithms 1 and 2 give the auditing gate and the end-to-end hardening procedure.

6 Experimental Setup

Corpus. Because sanctioned access to all six commercial APIs was not available at submission, we evaluate the pipeline on a *synthetic* corpus generated by `src/synthesize_demo.py`. It contains $N = 3600$ items over six placeholder generators (MODEL-A...MODEL-F), five subjects, and three difficulty levels ($6 \times 5 \times 3$ cells, 40 items per cell, $m = 4$ options each). Each placeholder generator has a *known*, *injected* length-bias strength ranging from strong (MODEL-A) to near-neutral (MODEL-F). **This corpus is not data from any real system**; its purpose is to (i) verify that the metrics recover known bias levels with correct ordering

and calibration, and (ii) demonstrate the mitigations on a controlled ground truth. Any resemblance between a placeholder’s value and a figure quoted for a named commercial model would be coincidental.

Procedure. `src/analyze.py` computes overall and per-group metrics with bootstrap CIs, fits the logistic leakage model, applies Holm–Bonferroni correction across per-model tests, runs the mitigation before/after comparison, and emits all tables and figures. The pipeline is deterministic given a seed.

Environment. Python 3.12 with NumPy, SciPy, pandas, Matplotlib, and statsmodels. All code, seeds, tables, and figures are in the accompanying repository.

7 Results

All results in this section are computed on the synthetic corpus of §6 and are illustrative of the pipeline, not measurements of real models.

7.1 Aggregate leakage

On the synthetic corpus, the always-longest strategy attains $LAA = 0.521$ (95% CI [0.506, 0.538]) against a chance baseline of 0.25, giving an Exploitability of $E = +0.361$ (95% CI [0.340, 0.382])—i.e. length alone carries a test-taker about 36% of the way from chance to a perfect cheat. The mean verbosity differential is $\bar{\delta} = +80.97$ characters (standardized effect 0.68; Wilcoxon $p \approx 2 \times 10^{-255}$). A length-only classifier achieves a leakage AUC = 0.698 (95% CI [0.688, 0.708]), and the logistic model gives an odds ratio of 2.25 per one standard deviation of option length (≈ 102 characters; $p \approx 0$). The key is the single longest option in 51.8% of items, versus the 25% expected under no signal (Fig. 2).

7.2 Variation across models

Table 1 reports per-generator metrics. By construction the injected bias produces a monotone gradient, which the metrics recover cleanly: Exploitability ranges from $E = +0.734$ (MODEL-A) down to $E = +0.109$ (MODEL-F), with leakage AUC tracking from 0.88 to 0.56. Every per-model test remains significant after Holm–Bonferroni correction. Figure 1 visualizes the gradient with confidence intervals. The point of this panel is methodological: it shows the suite is sensitive enough to rank generators by leakage and to separate a strongly leaking system from a near-neutral one—exactly the comparison a real cross-model study would need to make.

Table 1: Per-generator leakage on the **synthetic** corpus (illustrative). LAA: longest-answer accuracy; E : exploitability; $\bar{\delta}$: mean key–distractor character differential; AUC: length-only leakage AUC; p_{Holm} : Holm–Bonferroni-corrected p . “*” denotes $p_{\text{Holm}} < 0.001$.

Model	LAA	E	$\bar{\delta}$ (ch)	AUC	p_{Holm}
Model-A	0.801	+0.734	+172.4	0.883	6.2e-141 *
Model-B	0.643	+0.524	+122.6	0.794	3.0e-69 *
Model-C	0.533	+0.378	+84.3	0.719	1.4e-38 *
Model-D	0.417	+0.223	+47.1	0.628	4.8e-16 *
Model-E	0.399	+0.199	+37.3	0.604	2.5e-13 *
Model-F	0.332	+0.109	+22.2	0.563	2.3e-05 *




Figure 1: Exploitability E by generator on the synthetic corpus (illustrative). Bars show the normalized above-chance advantage of an always-longest strategy; whiskers are bootstrap 95% CIs. The gradient is injected ground truth that the metric recovers.

7.3 Variation across subject and difficulty

Figure 3 shows Exploitability across the subject $\times$ difficulty grid. Because the synthetic generator injects bias primarily as a function of model strength, the grid is comparatively flat across cells; in a real study this same heatmap is the instrument for detecting whether, for example, “hard” items leak more (a plausible hypothesis if models over-qualify difficult keys). We include it to demonstrate the per-cell estimator and its visualization.

7.4 Mitigation: before vs. after

We apply the hardening pipeline (Algorithm 2) to the synthetic corpus and recompute every metric. Table 2 and Figure 4 summarize the effect. The fraction of items flagged leaky falls from 71.4% to 2.9%. Exploitability drops from $E = +0.361$ to $E = -0.061$ —statistically indistinguishable from, and slightly below, chance—while the leakage AUC moves from 0.698 to 0.439 and the mean differential from +80.97 to −16.68 characters. In other words, after hardening, the longest option is no longer a reliable guess, and a length-only classifier does no better than chance. The mitigation over-corrects $\bar{\delta}$




Figure 2: Distribution of the correct option’s length *rank* (1 = longest) on the synthetic corpus (illustrative). Under no leakage every rank would hold 25%; the mass at rank 1 is the signature of answer-length leakage.




Figure 3: Exploitability E across subject $\times$ difficulty on the synthetic corpus (illustrative). The panel demonstrates the per-cell estimator that a real study would use to localize where leakage concentrates.

slightly negative; this is a tunable property of the deterministic matcher and can be centered by adjusting the target band.

7.5 Distributional view

Figure 5 shows the full distribution of the per-item verbosity differential before and after mitigation. Pre-mitigation mass sits to the right of zero (keys longer than distractors); post-mitigation the distribution re-centers near zero, consistent with the aggregate statistics above.

8 Discussion

Interpreting the synthetic results. The pipeline does what it is designed to do: the metrics recover injected bias with correct ordering and calibrated magnitudes, the statistical tests behave, and the mitigations remove the leak on controlled ground truth. This is necessary groundwork—an effect this consequential should

Table 2: Mitigation before/after on the synthetic corpus (illustrative). Lower magnitude is better for every row; the target for E , AUC, and $\bar{\delta}$ is parity (0, 0.5, 0).

Metric	Before	After	Δ
Longest-Answer Accuracy	0.521	0.204	-0.317
Exploitability E	0.361	-0.061	-0.422
Leakage AUC	0.698	0.439	-0.259
Mean differential $\bar{\delta}$ (chars)	80.969	-16.679	-97.648




Figure 4: Leakage metrics before and after hardening on the synthetic corpus (illustrative). The auditing gate plus length-balancing drives exploitability to parity and the length-only AUC to chance.

be measured with an instrument that has been validated end to end before it is pointed at production systems. What the synthetic study cannot tell us is the *real* magnitude for Claude, ChatGPT, Gemini, Grok, DeepSeek, or Perplexity; that requires running the harness against the live endpoints (§11).

Why a length signature is plausible in real models. Three non-exclusive mechanisms could produce keys that are systematically longer than distractors. (1) *Qualification necessity*: a fully correct statement often must include conditions and scope (“... under standard temperature and pressure”), whereas a plausible-but-wrong option can be crisp; correctness and length may be correlated for content reasons. (2) *Verbosity reward*: alignment from human feedback can reward thoroughness, and judges prefer longer answers [6, 7]; a model may carry this disposition into the option it is most committed to defending—the key. (3) *Distractor-generation strategy*: models may generate distractors by lightly perturbing or truncating the key, yielding shorter wrong options as a by-product. These mechanisms are hypotheses; behavioral testing with the harness can distinguish them (e.g., by checking whether the differential persists when the model is asked to justify each option), but confirming internal causes would require model-internal access we do not assume.




Figure 5: Per-item verbosity differential δ (key minus mean distractor, characters) before and after mitigation on the synthetic corpus (illustrative). Recentering near zero indicates the length signal has been removed.

Connection to established theory. Our metrics operationalize construct-irrelevant variance [2]: $E > 0$ is direct evidence that a non-content feature predicts the key, and a “pick-the-longest” test-taker is exercising classical test-wisdom [5]. The novelty is not the existence of the flaw, long known for human writers [3, 4], but its potential systematicity under LLMs and the fact that it can be *audited and repaired automatically* at generation time.

Implications for practice. The auditor (Algorithm 1) is cheap enough to run on every machine-authored quiz and to embed in authoring tools as a gate; it requires no model access, only the items. The hardening step integrates with any LLM authoring loop. For institutions, the form-level balancer offers a defense even when item-level differences are small. For researchers, the metric suite provides a common yardstick so that “which model leaks least” becomes an answerable, reproducible question.

9 Conclusion

We framed answer-length leakage in LLM-generated MCQs as a measurable threat to assessment validity, defined a reusable exploitability-oriented metric suite (LAA, E , $\bar{\delta}$, leakage AUC), built a provider-agnostic harness to collect items under a length-neutral prompt, and delivered a generation-time mitigation toolkit with a before/after protocol. On a clearly labeled synthetic corpus the suite recovers injected bias and the toolkit drives exploitability from $E = +0.36$ to parity while cutting leaky items from 71% to 3%. Answering the research questions: leakage *is* definable and quantifiable (RQ1); the suite resolves model-level differences (RQ2); a longest-option strategy beats chance on leaky corpora and length carries substantial information about

the key (RQ3); and lightweight mitigations remove the leak without discarding items (RQ4). The contribution is a deployable measurement- and mitigation stack; the open empirical question—the magnitude in real commercial systems—is one we have made *straightforward to answer* with the included harness and preregistered protocol.

10 Limitations

We are explicit about what this study does and does not establish.

- **No real-model measurements.** Every number in §7 comes from a synthetic corpus with injected bias. We do not report, and the reader must not infer, the leakage of any named commercial model. The synthetic study validates tooling only.
- **The reported magnitudes are illustrative.** The headline reduction ($E : +0.36 \rightarrow -0.06$) reflects controlled ground truth and tuned thresholds; real corpora may differ in both baseline leakage and achievable reduction.
- **Length is a proxy.** Character/word length is one surface cue; keys may also leak through specificity, grammaticality, or option position. We measure length; we do not claim it is the only or dominant cue.
- **Prompt sensitivity.** Real measurements will depend on the elicitation prompt, decoding temperature, and model version. Our length-neutral prompt is one defensible choice; results may not transfer across prompts.
- **Model drift.** Commercial models are updated frequently; any real-world estimate is a snapshot and may not persist.
- **Mechanism is hypothesized.** We offer behavioral hypotheses for why keys might be longer; we do not have model-internal access and do not claim to have identified the cause.
- **No human-subject data yet.** We have not measured how much a real “pick-the-longest” strategy raises scores for real students; that is a separate, IRB-governed study (§11).

11 Future Work and Recommendations

Preregistered real-model study. Run `src/collect.py` against the six assistants with institutional API access, holding the grid (5 subjects $\times$ 3 difficulties) and cell size fixed, with temperature and model versions logged. Pre-register the primary

hypothesis ($E > 0$ overall), the per-model comparison, and the subject/difficulty interaction. Estimate sample size for a target CI half-width on E ; the bootstrap machinery is already in place.

Human-subject validation of exploitability. With IRB approval and informed consent, present students with leaky and hardened forms and measure the score lift attributable to a longest-option strategy, and whether students spontaneously discover it. This converts E from a corpus statistic into a measured integrity risk.

Mechanism probes. Use the harness to test the qualification, verbosity-reward, and distractor-strategy hypotheses behaviorally (e.g., ask models to justify each option, or to generate distractors first), and study whether instruction-level interventions reduce leakage at the source.

Beyond length. Extend the metric suite to other construct- irrelevant cues (option specificity, “all/none of the above,” grammatical agreement, position) so that authoring tools can audit a battery of leaks, not just one.

Deployment. Package the auditor as a plug-in for common quiz and LMS authoring flows so that every machine-generated item is screened by default, turning the mitigation from a research artifact into infrastructure.

Reproducibility and Artifacts

All metrics (`metrics.py`), the collection harness (`collect.py`), the mitigation toolkit (`mitigate.py`), the synthetic generator (`synthesize_demo.py`), and the analysis driver (`analyze.py`) are included, along with the synthetic corpus, the generated tables, and all figures. The pipeline is deterministic given a seed. See Appendix A for the layout and commands.

Ethical Considerations

This work aims to *protect* assessment integrity. The metrics and harness could in principle help a student exploit leaky quizzes; we mitigate this by pairing every measurement tool with a remediation tool and by framing the contribution around auditing and repair. Any human-subject component must proceed under institutional review with informed consent and data minimization. We report on synthetic data precisely to avoid publishing unverified claims about named systems.

A Repository Layout and Commands

```
llm-mcq-length-bias/
  src/metrics.py      # LAA, E, delta, AUC, logistic
  src/collect.py      # 6-provider harness (--dry-run)
  src/mitigate.py     # auditor + 3 mitigations
  src/synthesize_demo.py # synthetic corpus generator
  src/analyze.py      # driver: tables + figures
  data/               # corpora (jsonl)
  results/            # results_*.json, *.tex tables
  figures/            # *.png/*.pdf
  paper/              # main.tex, references.bib
```

Typical workflow:

```
# 1. validate pipeline offline (no keys)
python src/collect.py --dry-run -o data/dry.jsonl
python src/synthesize_demo.py -o data/synthetic_mcq.jsonl
python src/analyze.py --corpus data/synthetic_mcq.jsonl \
    --tag synthetic

# 2. real run (needs your API keys, ToS-compliant)
export ANTHROPIC_API_KEY=... # etc. per provider
python src/collect.py --models all -o data/real.jsonl
python src/analyze.py --corpus data/real.jsonl --tag real
```

B Metric Summary

For a corpus with items $(s, \mathcal{O}, k)$ and length $\ell(\cdot)$: LAA is the longest-option score (Eq. 1); $E = (\text{LAA} - \bar{c}) / (1 - \bar{c})$ normalizes the above-chance advantage (Eq. 2); $\bar{\delta}$ is the mean key–distractor differential (Eq. 3); the leakage AUC is the key–outranks–distractor probability (Eq. 4); and the logistic odds ratio per SD expresses leakage as a single multiplier. All are reported with bootstrap CIs and per-group breakdowns.

References

- [1] Ghader Kurdi, Jared Leo, Bijan Parsia, Uli Sattler, and Salam Al-Emari. A systematic review of automatic question generation for educational purposes. *International Journal of Artificial Intelligence in Education*, 30(1):121–204, 2020.
- [2] Samuel Messick. Validity. In Robert L. Linn, editor, *Educational Measurement*, pages 13–103. American Council on Education / Macmillan, 3rd edition, 1989.
- [3] Thomas M. Haladyna, Steven M. Downing, and Michael C. Rodriguez. A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3):309–334, 2002.
- [4] Marie Tarrant, Aimee Knierim, Sasha K. Hayes, and James Ware. The frequency of item writing flaws in multiple-choice questions used in high stakes nursing assessments. *Nurse Education in Practice*, 6(6):354–363, 2006.

- [5] Jason Millman, Carol H. Bishop, and Robert Ebel. An analysis of test-wiseness. *Educational and Psychological Measurement*, 25(3):707–726, 1965.
- [6] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023.
- [7] Keita Saito, Akifumi Wachi, Koki Wataoka, and Youhei Akimoto. Verbosity bias in preference labeling by large language models, 2023. arXiv:2310.10076.
- [8] Ai-assisted MCQ creation increases item-writing flaws through automation bias: Evaluation of a workflow of multiple AI agents with teachers. *Frontiers in Computer Science*, 2026. <https://www.frontiersin.org/journals/computer-science/articles/10.3389/fcomp.2026.1831250/full>.
- [9] Thomas M. Haladyna and Steven M. Downing. A taxonomy of multiple-choice item-writing rules. *Applied Measurement in Education*, 2(1):37–50, 1989.
- [10] Steven M. Downing. The effects of violating standard item writing principles on tests and students: The consequences of using flawed test items on achievement examinations in medical education. *Advances in Health Sciences Education*, 10(2):133–143, 2005.
- [11] Michael C. Rodriguez. Three options are optimal for multiple-choice items: A meta-analysis of 80 years of research. *Educational Measurement: Issues and Practice*, 24(2):3–13, 2005.
- [12] Habibie Wicaksono Awalurahman and Indra Budi. Automatic distractor generation in multiple-choice questions: A systematic literature review. *PeerJ Computer Science*, 10:e2441, 2024.
- [13] Zhaopeng Qiu, Xian Wu, and Wei Fan. Automatic distractor generation for multiple choice questions in standard tests. In *Proceedings of the 28th International Conference on Computational Linguistics (COLING)*, pages 2096–2106, 2020.
- [14] Orchestrating LLM agents for scientific research: A pilot study of multiple choice question generation and evaluation. arXiv preprint, 2026. arXiv:2602.18891.
- [15] Comparative analysis of NLP-driven MCQ generators from text sources. *Computers and Education: Artificial Intelligence*, 2025. <https://www.sciencedirect.com/science/article/pii/S2666920X25000803>.
- [16] BenchMarker: An education-inspired toolkit for highlighting flaws in multiple-choice benchmarks. arXiv preprint, 2026. arXiv:2602.06221.
- [17] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations (ICLR)*, 2024.
- [18] Xinpeng Wang, Chengzhi Hu, Bolei Ma, Paul Röttger, and Barbara Plank. Look at the text: Instruction-tuned language models are more robust multiple choice selectors than you think, 2024. arXiv:2404.08382.